

AI Needs the Commons

Why Proprietary Practice Has Become Friction Rather Than Advantage

Abstract

This document argues that the productivity gain organisations seek from AI systems is conditional on a property those organisations do not control and often actively work against: the openness of the conventions, artefacts, and practices the work is built upon. AI systems are most fluent against densely represented public material — mainstream languages, open conventions, widely documented tools — and that fluency degrades in proportion to enclosure. Proprietary practice, historically a source of defensibility, therefore becomes an impedance against the most scalable input now available. The argument is not that AI ought to be open. It is that AI relies on a commons as a matter of operational fact, and that organisations wishing to exploit the productivity it promises must accept conventions this contradicts. No judgement is offered as to whether this is desirable. It is reported as a property of the technology.

1. A Plain Observation

There is a measurable difference between asking an AI system to work against a public, conventional environment and asking it to work against a private, proprietary one.

Point such a system at a mainstream language, a widely used version control platform, an open and well-documented framework, and the conventions it encounters are ones it has seen represented at enormous density. It requires little instruction. It produces output that fits the surrounding practice because the surrounding practice is the practice it is most fluent in.

Point the same system at a proprietary toolchain — an internal platform wrapped around a vendor-supplied system, an idiosyncratic workflow documented partly in restricted pages and partly in the heads of three long-serving engineers, conventions that exist as accidents of one organisation's history — and its marginal value falls. Not because the system is incapable, but because the thing it is fluent in is not present. It must be taught, at length, what it would otherwise have known by default.

This is not a complaint. It is a description. The fluency is a function of representational density, and representational density lives in the commons.

2. Where the Fluency Comes From

An AI system applied to knowledge work does not reason from first principles about an unfamiliar environment. It applies patterns at scale. The quality and cheapness of that application is governed by how densely the relevant patterns are represented in what the system has absorbed.

The relationship between fluency and density need not be asserted on its own. It follows from a property argued elsewhere: that the cost of changing a system is governed less by the state of its code than by *epistemic debt* — the

accumulated cost of understanding what a change will do. Epistemic debt accumulates wherever classification is implicit rather than formalised, wherever behaviour is relied upon without being observable, and wherever invariants exist only in human memory. Where it is high, change slows; where it is low, change remains cheap and safe.

An AI system is the actor for which epistemic debt is most absolute. It operates only on what is explicit and represented. It draws on none of the informal repair — the tacit memory, the inferred intent, the correction made by someone who will have to live with the result — by which humans continuously service the debt that their own practice accumulates. Its fluency is therefore not merely correlated with representational density; it is *governed* by epistemic debt directly. The system is fluent exactly where the debt is low and impotent exactly where it is high.

This reframes the commons precisely. Public conventions are densely represented because they are public — but more than that, they are *low epistemic debt by construction*. Open languages, shared tools, common frameworks, and the accumulated corpus of documentation and discussion that surrounds them are explicit, observable, and exercised in the open at enormous volume. The commons is the body of practice that has been made legible by the sheer fact of being shared.

Proprietary practice is, by construction, the opposite. It exists in one place, behind one boundary, exercised unobserved, described — where it is described

at all — in artefacts never intended to teach an outside reader anything, and held in large part as invariants in the memories of a few long-serving people. It is high epistemic debt in its purest form. Whatever advantage that condition once conferred on the humans who could service it, it subtracts directly from the fluency an AI system can bring to bear, because the AI cannot service it at all.

The substrate of fluency is the commons, and the reason is not mysterious. The commons is simply the part of practice whose epistemic debt has already been paid down by being made explicit and public. This is what the technology is.

3. The Operational Contradiction

The productivity gain that makes AI systems attractive to organisations is therefore conditional on a property those organisations' default instincts oppose.

Organisations differentiate. They enclose. They build proprietary practice precisely in order to be distinct, defensible, and hard to replicate. For much of the history of knowledge work this was rational: the bespoke internal system, the idiosyncratic convention, the privately accumulated way of doing things, constituted a moat.

Against an AI system, the same bespoke layer is no longer a moat. It is friction. Every deviation from public convention is a place the system is less fluent, a

place it must be expensively instructed, a place the promised productivity is not delivered.

So the organisation that wishes to capture the AI productivity surplus is asked to do something that contradicts its default mode: to align its practice with public convention rather than away from it; to prefer the densely represented over the privately distinct; to treat openness of practice not as a concession but as the precondition of the gain it is pursuing.

The surplus and the openness are not separable. The first is conditional on the second.

4. The Substrate Is Not Reproduced by Those Who Consume It

There is a second observation that follows from the first, and it is worth stating in operational terms.

The commons that produces the fluency was largely produced outside the organisations now extracting productivity from it. The open conventions, the public documentation, the shared corpus — these were generated by a distributed effort that the consuming organisation did not fund and does not maintain.

An organisation can take the fluency and contribute nothing back: keep its own practice private, its own conventions internal, its own artefacts enclosed. Nothing in the immediate transaction prevents this. But the fluency it

consumes depends on a substrate that such behaviour, generalised, does not replenish.

This has a practical consequence rather than a moral one. An organisation that severs itself from the commons — that insists on a wholly proprietary substrate — does not escape the dependency on representational density. It merely takes on the obligation to produce that density privately, at recurring cost, rebuilding behind its own boundary a thinner version of what the commons provides for free. Enclosure does not remove the precondition. It converts a shared and freely available substrate into a private one that must be perpetually and expensively reproduced.

The deeper cost is epistemic. To enclose is to choose, deliberately, to operate the AI against one's own highest-debt surface. The proprietary substrate is the part of practice whose epistemic debt has never been paid down — implicit, unobserved, held in memory — and it is precisely there that the AI is least able to help, because it is precisely there that there is nothing explicit for it to draw upon. The commons, low-debt and legible, sits available for nothing. The enclosed organisation declines it and turns the system loose on the one body of practice the system cannot read. The self-undermining is not only that the substrate goes unreplenished; it is that the firm has aimed its most scalable instrument at the surface where that instrument is structurally blind.

5. Enclosure Produces Ungrounded Output

The failure mode is not abstract. An AI system supplied only with thin, proprietary context produces output that is locally plausible and globally adrift — fluent in form, ungrounded in substance — because it has been cut off from the substrate that would otherwise ground it.

This is the same failure mode that arises whenever an AI system is given underspecified or impoverished input. Enclosure is one specific way of impoverishing it. The organisation that encloses everything, and feeds the system only its own thin internal context, receives precisely the fluent-but-ungrounded output that impoverished context produces.

It is worth being precise about what kind of failure this is, because it has a familiar shape. Declared architecture detaches from running behaviour when the record describes the system as it was imagined in meetings rather than as it exists in production; the document remains fluent while the reality it claims to describe drifts away beneath it. Enclosed AI output is the same detachment, generated rather than recorded. The proprietary context the system is fed is itself a declared description — thin, unobserved, unexercised — and the output faithfully realises *that description* while losing contact with the practice it was meant to serve. In both cases representation parts company with the reality it names, and in both cases the parting is invisible precisely because the representation remains internally coherent. Fluency is what detachment looks like from the inside.

The technical failure and the structural one are therefore the same failure. The output drifts because the grounding was withheld, and the grounding was withheld because the practice was enclosed.

6. The Honest Limit of the Argument

It does not follow that AI will force openness.

An organisation can refuse the commons. It can construct proprietary corpora, fine-tune against internal material, commission vendor-supplied systems pre-trained on private practice, and so manufacture a privately held substitute for the substrate the commons offers freely. This is possible. It is, in many domains, already happening.

The argument is therefore not that openness is inevitable, which would be a determinism worth distrusting. The argument is narrower and more durable: the easy, cheap, immediate productivity gain runs through the commons, and the enclosed alternative is available only at recurring expense — a treadmill of privately reproducing, in worse form, what the commons provides for nothing. Capital can pay to resist the commons. The point is that it must keep paying, and that the version it builds behind its own walls is thinner than the one it declined.

The contradiction does not resolve itself in openness' favour. It tilts the cost structure that way, and leaves the choice — and its price — to the organisation.

7. The Tacit and the Enclosed Are Not the Same Thing

A reader may object that the argument proves too much. If enclosed, tacit, unobserved practice is merely debt to be paid down, what becomes of the claim — argued elsewhere, and true — that human work depends on an irreducible substrate of tacit judgement and informal repair that *cannot* be made fully explicit without being destroyed? Is the tacit the precious thing that survives, or the debt to be eliminated? The argument here appears to want it both ways.

The objection is worth answering directly, because the resolution sharpens the whole position.

There are two distinct things being called tacit, and they must be separated. The first resists formalisation in principle. The second resists it only in fact.

The second is the easier to dispose of. It is practice that is unobserved not because observation would destroy it but because no one was obliged to make it legible: conventions that could have been written down and were not, invariants that happen to live in memory because habit never forced them out of it, classifications left implicit because the people who relied on them already shared them. This is not irreducible. It is contingent. It is epistemic debt in the ordinary sense, and it is debt precisely because it *could* be paid down and has not been.

The first is harder, and naming it correctly matters, because it is tempting to locate it in skill — in judgement too subtle to articulate, the expert's discrimination, the craft that resists being written. That is real, but it is not the

floor. The floor is lower and more structural than skill. What the AI system irreducibly lacks is not refined judgement but *position*: it is unsituated, unexposed, and indifferent to outcome. It occupies no body, bears no consequence, and stands nowhere in particular. The irreducible human remainder is not, at bottom, that the human knows things too subtle to formalise; it is that the human is *exposed* — finite, situated, and liable to bear the cost of being wrong. Judgement that cannot be formalised is largely downstream of this: it is the form taken by a situated and vulnerable agent's ongoing accommodation to a world it cannot survey from outside. The repair the AI cannot perform is the repair of one who has something at stake.

This distinction does the work. Enclosure does not protect the first kind and could not, because the first kind was never available for externalisation to anything unembodied, open practice or closed. Enclosure merely accumulates the second kind — contingent debt — and then feeds it to a system that cannot service it, and the output drifts. The genuinely irreducible remainder stays exactly where it always was: not in the corpus, open or proprietary, but in the embodied party who is answerable for what the system produces because they are the one the consequences fall upon.

The commons, then, is not the elimination of the tacit, and it makes no claim on the irreducible. It is simply the part of contingently tacit practice that the discipline of being shared has already made explicit. What remains irreducible — the situated, exposed, answerable position from which alone the output can be judged worth trusting — remains so, in the commons as everywhere else.

The error of enclosure is not that it guards something precious. It guards nothing precious; it merely fails to externalise the contingent kind while leaving the irreducible kind precisely where it would have been anyway, and in doing so confuses the two — treating ordinary unpaid epistemic debt as though it were the inviolable human remainder, and thereby declining to pay down a cost the technology has made newly expensive to carry.

8. Practice

This is not a speculative claim. A development framework built around an AI coding agent works well to the precise extent that it leans on public convention: a mainstream language, an open version control platform, widely documented testing and specification practices, a stack the agent has encountered at density. The same framework applied to a wholly proprietary substrate — a closed platform, a vendor-locked system, conventions documented nowhere public — would have demanded far more instruction to achieve far less, and would have been perpetually vulnerable to the ungrounded-output failure mode for want of a grounding the commons would otherwise have supplied.

The discipline that makes such a framework effective is, in large part, the discipline of staying close to public convention wherever the work permits, and of being explicit — at cost — wherever it does not. That second clause is the paying-down of contingent epistemic debt, performed deliberately, precisely where the commons does not reach.

9. Conclusion

AI systems do not have a politics. But they have a precondition, and the precondition is a commons.

Their fluency is a function of how densely the relevant practice is represented in shared, public material. Enclosure subtracts from that density. Proprietary practice, once a source of defensibility, is now an impedance against the most scalable input available — and the productivity surplus an organisation seeks is conditional on an openness its default instincts oppose.

Organisations are free to enclose. They are free to rebuild the substrate privately, at recurring cost, and to accept the thinner fluency that results. Nothing here recommends against it.

It is offered only as a fact of what the technology is: proprietary thinking has become counterproductive to the gain it was meant to protect. What form an organisation gives to that fact — whether it opens its practice to exploit the surplus, or pays to enclose and reproduces a lesser substrate behind its walls — is for the organisation to determine.

The technology relies on a commons. The rest is a choice, and its price.