

Does AI Need the Commons?

On Legibility, Enclosure, and the Cost of Being Read

0. Frame

There is a familiar observation that AI systems work better against open, public, conventional practice than against private, proprietary practice. The observation is usually offered as a technological fact: a property of the tools, reported neutrally, from which organisations may draw whatever operational conclusions they wish.

This paper takes the observation seriously and the neutrality less so. The question of whether AI depends on a commons is not a purely technological one, and any answer that presents itself as purely technological has already taken a position — because to report an arrangement of cost and access as a property of the technology is to naturalise it, and naturalising a contested arrangement is not a neutral act. The question is at once technological, economic, and political, and the three cannot be cleanly separated without the separation itself doing work that ought to be argued rather than assumed.

The aim here is therefore not to assert that AI needs the commons, nor to deny it, but to ask the question properly and follow it to where it actually resolves. It resolves, as will be argued, into an economic and political claim rather than a technological one — and the path by which a seemingly technological question

arrives at a political answer is itself the most important thing the paper has to report.

1. The Observation, as Observation

Point an AI coding system at a mainstream language, a widely used version control platform, an open and well-documented framework, and it requires little instruction. It produces output that fits the surrounding practice, because the surrounding practice is the practice it is most fluent in. The conventions it meets are ones it has encountered at enormous density.

Point the same system at a proprietary toolchain — an internal platform wrapped around a vendor-supplied system, a workflow documented partly in restricted pages and partly in the memory of a few long-serving engineers, conventions that exist as accidents of one organisation's history — and its marginal value falls. Not because the system is incapable, but because the thing it is fluent in is not present. It must be taught, at length, what it would otherwise have known by default.

This is a real and currently observable difference. It is the thing to be explained. It is not yet an explanation. The temptation is to explain it by appeal to *publicness* — the system is better against public practice because the practice is public — and to stop there. The remainder of this paper is an argument that publicness is not the operative variable, that it correlates with the operative variable for a reason worth stating, and that once the real variable is named the question changes character.

2. Disaggregating the Variable

The candidate explanation, “publicness,” contains two distinct properties that ordinarily travel together and need to be pulled apart.

The first is **representational density**: the system has encountered the relevant patterns at volume, and applies them cheaply because it has seen them many times. The second is **legibility**: the practice is explicit, observable, and exercised in a form that can be read off rather than inferred. Public conventions tend to score high on both, which is why publicness appears to be the explanation. But density and legibility are separable, and separating them is the move the rest of the argument runs on.

It is worth asking why the two have historically travelled together, because the answer supplies the mechanism that a bare appeal to publicness leaves implicit. Practice that is exercised in the open, at volume, by many parties who did not coordinate in advance, is placed under continuous pressure to become legible. People who must use a convention they did not author have to be able to read it; disagreements have to be settled in terms others can follow; the sheer fact of shared use forces what would otherwise remain implicit to be written down, named, and made inspectable. Legibility, in the commons, is not a virtue deliberately pursued. It is a by-product of practice being shared under conditions where no single party can hold the whole thing in memory. Publicness produces legibility because publicness removes the option of relying on private, unarticulated understanding.

This reframes the original observation. The system is fluent against the commons not because the commons is *public* but because the commons is *legible*, and it is legible because being shared at volume is the condition under which contingent, unwritten understanding gets forced into explicit form. Publicness is the historical cause of legibility, not its substance. The substance is legibility itself — and legibility, unlike publicness, is a property an organisation can in principle produce by other means.

3. The Commons Is Neither Sufficient Nor Necessary

If the operative variable is legibility rather than publicness, two consequences follow that the simple version of the thesis cannot accommodate. Both are observable; neither is hypothetical.

The commons is not sufficient. The public corpus is dense but not uniformly legible. It contains contradictory conventions, deprecated practice presented alongside current practice with no marker distinguishing them, and large volumes of low-quality material optimised for visibility rather than correctness. Density without legibility does not produce grounded output; it produces output that is fluent in form and wrong in substance, because the system reproduces a pattern it has seen at volume without the means to know the pattern is stale or contested. A high-density, low-legibility region of the commons is, for practical purposes, one of the *least* reliable environments to operate in, precisely because the fluency is convincing and the grounding is absent. Density is not the good. Legible density is.

The commons is not necessary. A proprietary environment that has been made deliberately legible — clean documentation, enforced conventions, observable behaviour, tests that exercise the system’s actual invariants — can present an agent with a more readable surface than mainstream-but-chaotic public practice. Such environments are not the norm; proprietary practice is, in the ordinary case, *less* legible than open practice, not more. But the reason is instructive. Enclosure does not cause illegibility directly. It removes the pressure that produces legibility for free. In the open, legibility is forced by shared use; behind a boundary, nothing forces it, so its default state is to remain unwritten — held in memory, relied upon without being recorded. Illegibility is the *default* outcome of enclosure, not its necessary consequence. An organisation can pay to produce legibility privately. It simply must pay, because the thing that produces it for nothing in the commons — shared use under conditions that forbid private understanding — is exactly the thing enclosure removes.

The net finding of this section: AI does not need the *commons*. It needs *legibility*. The commons is the cheapest and most scalable source of legibility currently available, but it is one source among possible others, and its advantage is economic rather than technological — it supplies legibility that someone else has already paid to produce.

4. The Floor Beneath Legibility

Before drawing the economic consequence, one distinction has to be made, because without it the argument appears to prove too much. If legibility is

simply unwritten understanding waiting to be written down, then everything tacit is merely debt to be paid, and the familiar claim that human work rests on an irreducible substrate of judgement that cannot be fully formalised would seem to be denied. It is not denied. But two different things are being called tacit, and they behave differently.

The first kind resists formalisation only *in fact*. It is unwritten not because writing it would destroy it but because no one was obliged to make it legible: conventions that could have been recorded and were not, invariants that live in memory because habit never forced them out, classifications left implicit because the people who relied on them already shared them. This is contingent. It is legibility deferred, and it can be produced — at cost — whenever the deferral stops being affordable. It is exactly the kind of understanding that shared use forces into the open in the commons, and that enclosure leaves unforced.

The second kind resists formalisation *in principle*, and it is tempting to locate it in skill — judgement too subtle to articulate, the expert's discrimination, the craft that will not be written down. That is real, but it is not the floor. The floor is lower and more structural than skill. What an AI system irreducibly lacks is not refined judgement but *position*. It is unsituated, exposed to no consequence, and indifferent to outcome. It occupies no particular place, bears no cost when it is wrong, and stands nowhere that the result can fall back upon. Judgement that cannot be formalised is largely downstream of this: it is the form taken by a situated and answerable agent's ongoing accommodation to a world it cannot

survey from outside. The remainder that no corpus contains is not, at bottom, subtle knowledge. It is the fact of being the party to whom the consequences accrue.

This distinction does the work the evaluation needs. A deliberately legible proprietary environment can pay down the *first* kind of tacit understanding behind its boundary, which is precisely why enclosure is not technologically fatal — the contingent debt can be settled privately, at a price. But no environment, open or closed, touches the *second* kind, because the second kind was never available for externalisation to anything unsituated, in the commons or anywhere else. It remains where it always was: with the embodied party who is answerable for what the system produces, because they are the one the consequences fall upon. The commons makes no claim on the irreducible remainder. It is simply the portion of *contingently* tacit practice that the discipline of being shared has already made explicit.

A note on a failure mode that follows from the first kind. When a system is fed a thin, unexercised description of a practice rather than the practice as it is actually performed, it produces output that realises the description while losing contact with what the description was meant to name — fluent in form, adrift in substance. This is the same shape as a more familiar failure: declared architecture detaching from running behaviour, the record describing the system as it was imagined rather than as it runs. In both cases representation parts company with the reality it names, and in both cases the parting is hard to see because the representation remains internally coherent. Fluency is what

that detachment looks like from the inside. The two failures are not identical — they have different causes and different remedies — but they share a shape, and the shape is the same one this whole argument turns on: the gap between what is written and what is the case, which legibility narrows and enclosure, left to its default, widens.

5. The Two Ledgers

Once the operative variable is legibility, and legibility is something the commons supplies cheaply because someone else has already produced it, the question stops being technological and becomes a question about cost — specifically, about who bears it and on whose account it appears.

An organisation drawing fluency from the commons is drawing on legibility it did not produce. The open conventions, the public documentation, the shared corpus that make an agent fluent were generated by a distributed effort the organisation did not fund and does not maintain. It can consume that legibility and contribute nothing toward it: keep its own practice private, its own conventions internal, its own artefacts unshared. Nothing in the immediate transaction prevents this, and on the organisation's own books the arrangement is simply efficient — it receives a productive input at no cost.

This is where the level at which the cost is counted becomes decisive, because the same arrangement reads completely differently depending on the ledger it is entered on. On the organisation's ledger, consuming the commons while contributing nothing is free. On the ledger of the substrate as a whole, the

same behaviour, if general, does not replenish what it draws on: the legibility the commons supplies has to be produced by someone, and a population of pure consumers produces none of it. The cost has not vanished. It has been moved to an account no single organisation is charged against, and therefore one that no rational organisation, optimising its own ledger, has reason to see.

It is worth being exact about what the commons actually supplies, because the ordinary economic description undersells it. The familiar reading is that the commons provides an underpriced input, and that the remedy for an underpriced input is to price it. But the legibility the commons produces is not simply cheap labour that could be costed and billed. It is produced under a different arrangement altogether — distributed, uncoordinated, performed by parties for reasons that are frequently not commercial, under conditions where legibility emerges as a by-product of shared use rather than as a deliverable anyone was contracted to produce. The significant fact about enclosure is not that it underpays for this. It is that it cannot reproduce the arrangement that generates it. An organisation that encloses does not pay a fair price for the legibility it used to get free; it takes on the obligation to produce that legibility under conditions that are structurally worse at producing it — coordinated where the commons was distributed, contracted where the commons was incidental, bounded where the commons was open — and obtains a thinner result at recurring cost. The commons is not cheap legibility. It is legibility produced by a mechanism that enclosure does not have and cannot easily build.

Neither ledger is the wrong one. That is the point. The organisation's ledger is correct about the organisation, and on it enclosure can be a defensible choice with a calculable price. The wider ledger is correct about the substrate, and on it a population of enclosers is drawing down a resource none of them maintains. The two are both accurate and they do not agree, and the disagreement is not a problem to be resolved by finding the true cost. The disagreement is the finding. A question that presented itself as technological — does the tool need the commons — has turned into a question about which account the cost is entered on, and that is not a question the technology can answer.

6. Where the Question Resolves

Does AI need the commons? Technologically, no — and the “no” is the result, not an evasion of one. AI needs legibility. The commons is the cheapest, most scalable, and externally produced source of legibility available, which is why the easy and immediate productivity runs through it and why staying close to public convention, wherever the work permits, is the path of least resistance. But the dependency is on legibility, not on the commons as such. An organisation can produce legibility privately. It can construct internal corpora, document its proprietary practice to a high standard, and so manufacture behind its own boundary a substitute for what the commons supplies. This is possible, and in some domains already routine. The technological question closes, and it closes against necessity.

What it closes into is not nothing. It is a cost structure with two properties. The first is that the private substitute is available only at recurring expense, and is in the ordinary case thinner than the commons it replaces, because the mechanism that produces legibility cheaply — shared use under conditions that forbid private understanding — is the precise mechanism enclosure removes. The second, and the one the technological framing is built to obscure, is that the cost of the commons does not disappear when an organisation consumes it for free. It is borne elsewhere, on an account the organisation has no reason to read, by a mechanism that pure consumption does not sustain.

This is why the question does not resolve technologically and should not be made to. The apparent technological necessity — that AI depends on open practice — was never a property of the tool. It was a fact about where legibility is currently cheapest to obtain, and where it is cheapest to obtain is a question about who has produced it and who is charged for producing it. The technological analysis does not hand the question off to economics and politics once the real work is done. It reveals that the question was economic and political from the start, wearing the costume of a tool's capability. The variable everyone took to be technical — public practice yields fluent output — was a distribution of cost the whole time.

AI does not need the commons. It needs what the commons has, until now, produced more cheaply than anything else: a body of practice made legible by the conditions of being shared. Whether that remains the cheapest source, whether the private substitutes improve, whether the substrate continues to be

produced by those who increasingly only consume it — these are not settled by the capabilities of the technology. They are settled by choices about where the cost of legibility is to fall, made by parties each reading a ledger that is accurate about itself and silent about the rest. The technology has no view on the matter. The arrangement does, and the arrangement is the thing worth seeing clearly — which is the one thing a strictly technological account of it is constructed not to show.